

Research Interests

Your Name

2026-06-27

Write your post here.

Ideas

1. SDRM
2. SiCPO
3. RL on Cartridges
4. **SDRM — Self-Distilled Reward Modeling:** Uses a privileged version of the model to evaluate and compare its own on-policy outputs, then distills those judgments into a reusable reward model for RL—learning from self-evaluation rather than imitating teacher responses.
5. **SiCPO — Self-Informed Critic Policy Optimization:** Uses privileged post-outcome information to train a critic that assigns dense, step-level credit across a reasoning trajectory, identifying which intermediate decisions increased or decreased the probability of eventual success. Like an implicit value function with TD bootstrapping/Generalized Advantage Estimation (GAE).
6. **RL on Cartridges:** Freezes the base model and applies reinforcement learning directly to a small customer- or task-specific memory module, such as a learned KV-cache cartridge, so it acquires specialized agent skills without modifying the full model.
 - SDRM: Can a model’s privileged self-knowledge be converted into a reliable reward signal without copying its answers?
 - SiCPO: Can hindsight about final outcomes be used to assign accurate credit to individual reasoning decisions?
 - CARL: Can useful agent behavior be learned through RL in a modular experience layer while leaving the base model frozen?